

ModelTune_LogReg-XGBoost_ComboPlots

October 1, 2022

```
[ ]: !pip install xgboost
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-  
wheels/public/simple/  
Requirement already satisfied: xgboost in /usr/local/lib/python3.7/dist-packages  
(0.90)  
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages  
(from xgboost) (1.21.6)  
Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages  
(from xgboost) (1.7.3)
```

```
[2]: %cd /content/drive/MyDrive/Github/ml-blog
```

```
/content/drive/MyDrive/Github/ml-blog
```

```
[3]: import os  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
from xgboost import XGBClassifier  
from sklearn.linear_model import LogisticRegression  
from sklearn.model_selection import train_test_split, GridSearchCV  
from sklearn.metrics import confusion_matrix, classification_report
```

0.1 Import Scaled Data

```
[ ]: filename = 'taiwan-credit-col-transform-FULL.csv'
```

```
[ ]: DATA = os.path.relpath('/content/drive/MyDrive/Github/ml-blog/credit/data/' +  
    ↪filename)  
  
df = pd.read_csv(DATA, index_col='ID', header=0)  
df.head()
```

```
[ ]:      category__SEX_1  category__SEX_2  category__EDUCATION_0  \  
ID  
1          0.0          1.0          0.0  
2          0.0          1.0          0.0
```

3	0.0	1.0	0.0
4	0.0	1.0	0.0
5	1.0	0.0	0.0

	category__EDUCATION_1	category__EDUCATION_2	category__EDUCATION_3	\
ID				
1	0.0	1.0	0.0	
2	0.0	1.0	0.0	
3	0.0	1.0	0.0	
4	0.0	1.0	0.0	
5	0.0	1.0	0.0	

	category__EDUCATION_4	category__EDUCATION_5	category__EDUCATION_6	\
ID				
1	0.0	0.0	0.0	
2	0.0	0.0	0.0	
3	0.0	0.0	0.0	
4	0.0	0.0	0.0	
5	0.0	0.0	0.0	

	category__MARRIAGE_0	...	nums__BILL_AMT4	nums__BILL_AMT5	\
ID		...			
1	0.0	...	-0.672497	-0.663059	
2	0.0	...	-0.621636	-0.606229	
3	0.0	...	-0.449730	-0.417188	
4	0.0	...	-0.232373	-0.186729	
5	0.0	...	-0.346997	-0.348137	

	nums__BILL_AMT6	nums__PAY_AMT1	nums__PAY_AMT2	nums__PAY_AMT3	\
ID					
1	-0.652724	-0.341942	-0.227086	-0.296801	
2	-0.597966	-0.341942	-0.213588	-0.240005	
3	-0.391630	-0.250292	-0.191887	-0.240005	
4	-0.156579	-0.221191	-0.169361	-0.228645	
5	-0.331482	-0.221191	1.335034	0.271165	

	nums__PAY_AMT4	nums__PAY_AMT5	nums__PAY_AMT6	default payment next month
ID				
1	-0.308063	-0.314136	-0.293382	1
2	-0.244230	-0.314136	-0.180878	1
3	-0.244230	-0.248683	-0.012122	0
4	-0.237846	-0.244166	-0.237130	0
5	0.266434	-0.269039	-0.255187	0

[5 rows x 92 columns]

```
[ ]: X = df.iloc[:, :91]
      y = df.iloc[:, 91]
      X.columns, y.name, y.shape
```

```
[ ]: (Index(['category__SEX_1', 'category__SEX_2', 'category__EDUCATION_0',
            'category__EDUCATION_1', 'category__EDUCATION_2',
            'category__EDUCATION_3', 'category__EDUCATION_4',
            'category__EDUCATION_5', 'category__EDUCATION_6',
            'category__MARRIAGE_0', 'category__MARRIAGE_1', 'category__MARRIAGE_2',
            'category__MARRIAGE_3', 'category__PAY_0_-2', 'category__PAY_0_-1',
            'category__PAY_0_0', 'category__PAY_0_1', 'category__PAY_0_2',
            'category__PAY_0_3', 'category__PAY_0_4', 'category__PAY_0_5',
            'category__PAY_0_6', 'category__PAY_0_7', 'category__PAY_0_8',
            'category__PAY_2_-2', 'category__PAY_2_-1', 'category__PAY_2_0',
            'category__PAY_2_1', 'category__PAY_2_2', 'category__PAY_2_3',
            'category__PAY_2_4', 'category__PAY_2_5', 'category__PAY_2_6',
            'category__PAY_2_7', 'category__PAY_2_8', 'category__PAY_3_-2',
            'category__PAY_3_-1', 'category__PAY_3_0', 'category__PAY_3_1',
            'category__PAY_3_2', 'category__PAY_3_3', 'category__PAY_3_4',
            'category__PAY_3_5', 'category__PAY_3_6', 'category__PAY_3_7',
            'category__PAY_3_8', 'category__PAY_4_-2', 'category__PAY_4_-1',
            'category__PAY_4_0', 'category__PAY_4_1', 'category__PAY_4_2',
            'category__PAY_4_3', 'category__PAY_4_4', 'category__PAY_4_5',
            'category__PAY_4_6', 'category__PAY_4_7', 'category__PAY_4_8',
            'category__PAY_5_-2', 'category__PAY_5_-1', 'category__PAY_5_0',
            'category__PAY_5_2', 'category__PAY_5_3', 'category__PAY_5_4',
            'category__PAY_5_5', 'category__PAY_5_6', 'category__PAY_5_7',
            'category__PAY_5_8', 'category__PAY_6_-2', 'category__PAY_6_-1',
            'category__PAY_6_0', 'category__PAY_6_2', 'category__PAY_6_3',
            'category__PAY_6_4', 'category__PAY_6_5', 'category__PAY_6_6',
            'category__PAY_6_7', 'category__PAY_6_8', 'nums__LIMIT_BAL',
            'nums__AGE', 'nums__BILL_AMT1', 'nums__BILL_AMT2', 'nums__BILL_AMT3',
            'nums__BILL_AMT4', 'nums__BILL_AMT5', 'nums__BILL_AMT6',
            'nums__PAY_AMT1', 'nums__PAY_AMT2', 'nums__PAY_AMT3', 'nums__PAY_AMT4',
            'nums__PAY_AMT5', 'nums__PAY_AMT6'],
           dtype='object'), 'default payment next month', (30000,))
```

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
      ↪random_state=42)
```

0.2 Logistic Regression Baseline

```
[ ]: lr_base_clf = LogisticRegression(penalty='l2', tol=0.0001,
      C=1, solver='lbfgs',
      max_iter=500, multi_class='ovr').fit(X_train,
      ↪y_train)
```


```
[ ]: y_lrbase_preds = lr_base_clf.predict(X_test)
```

```
[ ]: print(confusion_matrix(y_test, y_lrbase_preds))
```

```
[[4461  226]
 [ 861  452]]
```

```
[ ]: sns.heatmap(confusion_matrix(y_test, y_lrbase_preds), cmap='coolwarm',
↪annot=True, fmt='g')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10f0898b50>
```



```
[ ]: print(classification_report(y_test, y_lrbase_preds))
```

	precision	recall	f1-score	support
0	0.84	0.95	0.89	4687
1	0.67	0.34	0.45	1313
accuracy			0.82	6000
macro avg	0.75	0.65	0.67	6000
weighted avg	0.80	0.82	0.80	6000

It's about what we expected....

0.3 GridSearchCV to compare hyperparameter settings

```
[ ]: lr_param_grid = [
    {'penalty': ['l2'], 'C': [0.1, 1.0, 5.0, 10.0, 100.0], 'tol': [0.001, 0.01, 0.5, 1.0]},
    {'penalty': ['l1'], 'C': [0.1, 1.0, 5.0, 10.0, 100.0], 'tol': [0.001, 0.01, 0.5, 1.0]},
    {'penalty': ['none'], 'tol': [0.001, 0.01, 0.5, 1.0]},
]

xgb_param_grid = [
    {'max_depth': [2], 'n_estimators': [25, 50, 100, 500], 'learning_rate': [0.01, 0.05, 0.075, 0.1]},
    {'max_depth': [4], 'n_estimators': [25, 50, 100, 500], 'learning_rate': [0.01, 0.05, 0.075, 0.1]},
    {'max_depth': [6], 'n_estimators': [25, 50, 100, 500], 'learning_rate': [0.01, 0.05, 0.075, 0.1]},
]

lr = LogisticRegression(max_iter=2000, solver='saga')

xgb = XGBClassifier(use_label_encoder=False)
```

```
[ ]: lr_clf = GridSearchCV(lr, param_grid=lr_param_grid, scoring='roc_auc')
lr_clf.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(estimator=LogisticRegression(max_iter=2000, solver='saga'),
                  param_grid=[{'C': [0.1, 1.0, 5.0, 10.0, 100.0], 'penalty': ['l2'],
                                'tol': [0.001, 0.01, 0.5, 1.0]},
                                {'C': [0.1, 1.0, 5.0, 10.0, 100.0], 'penalty': ['l1'],
                                'tol': [0.001, 0.01, 0.5, 1.0]},
                                {'penalty': ['none'], 'tol': [0.001, 0.01, 0.5, 1.0]}],
                  scoring='roc_auc')
```

```
[ ]: print(lr_clf.best_params_, lr_clf.best_estimator_, round(lr_clf.best_score_, 3))

{'C': 0.1, 'penalty': 'l2', 'tol': 0.001} LogisticRegression(C=0.1,
max_iter=2000, solver='saga', tol=0.001) 0.768
```

```
[ ]: xgb_clf = GridSearchCV(xgb, param_grid=xgb_param_grid, scoring='roc_auc')
xgb_clf.fit(X_train, y_train)
```

```
[ ]: GridSearchCV(estimator=XGBClassifier(use_label_encoder=False),
                  param_grid=[{'learning_rate': [0.01, 0.05, 0.075, 0.1],
                                'max_depth': [2],
                                'n_estimators': [25, 50, 100, 500]},
                                {'learning_rate': [0.01, 0.05, 0.075, 0.1],
                                'max_depth': [4],
```

```

        'n_estimators': [25, 50, 100, 500]],
        {'learning_rate': [0.01, 0.05, 0.075, 0.1],
         'max_depth': [6],
         'n_estimators': [25, 50, 100, 500]]},
        scoring='roc_auc')

```

```

[ ]: print(xgb_clf.best_params_, xgb_clf.best_estimator_, round(xgb_clf.best_score_,
↪3))

```

```

{'learning_rate': 0.05, 'max_depth': 6, 'n_estimators': 100}
XGBClassifier(learning_rate=0.05, max_depth=6, use_label_encoder=False) 0.783

```

0.4 Evaluation, Graphs, etc.

```

[ ]: from sklearn.metrics import roc_curve, auc, roc_auc_score

```

```

[ ]: y_lr_probas = lr_clf.predict_proba(X_test)[:, 1] # using the 'greater label', i.
↪e., 0 for 'No default'.

```

```

[ ]: y_xgb_probas = xgb_clf.predict_proba(X_test)[:, 1]

```

```

[ ]: y_xgb_probas.shape

```

```

[ ]: (6000,)

```

```

[ ]: y_test_np = y_test.to_numpy()

```

```

[ ]: y_test_np.shape

```

```

[ ]: (6000,)

```

```

[ ]: truePR_lr = dict()
    falsePR_lr = dict()
    roc_auc_lr = dict()

    truePR_xgb = dict()
    falsePR_xgb = dict()
    roc_auc_xgb = dict()

    falsePR_lr, truePR_lr, _ = roc_curve(y_test_np, y_lr_probas)
    roc_auc_lr = auc(falsePR_lr, truePR_lr)

    falsePR_xgb, truePR_xgb, _ = roc_curve(y_test_np, y_xgb_probas)
    roc_auc_xgb = auc(falsePR_xgb, truePR_xgb)

    plt.style.use('ggplot')

```

```


plt.figure(figsize=(16,12))
lw = 2

plt.plot(
    falsePR_lr,
    truePR_lr,
    color='darkorange',
    lw=lw,
    label="ROC Curve LR Model (area = %0.3f)" % roc_auc_lr,
)

plt.plot(
    falsePR_xgb,
    truePR_xgb,
    color='red',
    linestyle='dotted',
    lw=lw,
    label="ROC Curve XGB Model (area = %0.3f)" % roc_auc_xgb,
)

plt.plot([0,1], color='navy', lw=lw, linestyle='--', label='Baseline 0.5 - Lack
↳of Skill')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) on LogisticRegression and
↳XGBoostClassifier')
plt.legend(loc='lower right')
plt.savefig('./credit/model/best-combo-model-auc.png')
plt.show()

```



```
[ ]: y_xgb_probas = xgb_clf.predict_proba(X_test)[: , 1] # using the 'greater label',  
      ↪ i.e., 0 for 'No default'.
```

```
[ ]: truePR = dict()  
      falsePR = dict()  
      roc_auc = dict()  
  
      falsePR, truePR, _ = roc_curve(y_test_np, y_xgb_probas)  
      roc_auc = auc(falsePR, truePR)  
  
      plt.style.use('ggplot')  
      plt.figure(figsize=(16,12))  
      lw = 2  
      plt.plot(  
          falsePR,  
          truePR,  
          color='red',  
          lw=lw,  
          label="ROC Curve (area = %0.3f)" % roc_auc,
```




```

)

plt.plot([0,1], color='navy', lw=lw, linestyle='--', label='Baseline 0.5 - Lack
↳of Skill')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) on
↳XGBClassifier(learning_rate=0.05, max_depth=6, use_label_encoder=False)')
plt.legend(loc='lower right')
plt.savefig('./credit/model/best-xgb-model-auc.png')
plt.show()

```



```

[ ]: y_lr_preds = lr_clf.predict(X_test)
     y_xgb_preds = xgb_clf.predict(X_test)

```

```

[ ]: print(classification_report(y_test_np, y_lr_preds))

```

```


precision    recall  f1-score   support

```

0	0.84	0.95	0.89	4687
1	0.66	0.35	0.46	1313
accuracy			0.82	6000
macro avg	0.75	0.65	0.67	6000
weighted avg	0.80	0.82	0.80	6000

```
[ ]: sns.heatmap(confusion_matrix(y_test, y_lr_preds), cmap='coolwarm', annot=True,
    ↪fmt='g')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10ecd18c50>
```




```
[ ]: print(classification_report(y_test_np, y_xgb_preds))
```

	precision	recall	f1-score	support
0	0.84	0.95	0.89	4687
1	0.68	0.35	0.47	1313
accuracy			0.82	6000
macro avg	0.76	0.65	0.68	6000
weighted avg	0.81	0.82	0.80	6000

```
[ ]: sns.heatmap(confusion_matrix(y_test, y_xgb_preds), cmap='coolwarm', annot=True,
    ↪fmt='g')
```

```
[ ]: <matplotlib.axes._subplots.AxesSubplot at 0x7f10eb332210>
```



0.5 SMOTE to Address Class Imbalance

- The dataset has a clear imbalance in the label counts. This means our models will have trouble discerning a **minority class** instance clearly.
- We can use the **SMOTE** method to improve class balance.

```
[ ]: from imblearn.under_sampling import RandomUnderSampler
    from imblearn.over_sampling import SMOTE
    from imblearn.pipeline import Pipeline
    from sklearn.metrics import classification_report
```

```
[ ]: ov = SMOTE(sampling_strategy=0.5, random_state=42)
    und = RandomUnderSampler(sampling_strategy=0.35, random_state=42)
```

```
[ ]: from collections import Counter
```

```
[ ]: count = Counter(y)
    print(count, 'Ratio: ', round(count[1]/count[0], 4), '% minority class (1)')
```

```
Counter({0: 23364, 1: 6636}) Ratio: 0.284 % minority class (1)
```

```
[ ]: count_train = Counter(y_train)
print(count_train, 'Ratio: ', round(count_train[1]/count_train[0], 4), '% minority class (1)')
```

Counter({0: 18677, 1: 5323}) Ratio: 0.285 % minority class (1)

```
[ ]: test_pipe = Pipeline(steps=[('und', und), ('ov', ov)]))
```

```
[ ]: X_res, y_res = test_pipe.fit_resample(X_train, y_train)
```

```
[ ]: count_res = Counter(y_res)
print(count_res, 'Ratio: ', round(count_res[1]/count_res[0], 4), '% minority class (1)')
```

Counter({0: 15208, 1: 7604}) Ratio: 0.5 % minority class (1)

```
[ ]: # {'learning_rate': 0.05, 'max_depth': 6, 'n_estimators': 100}
      ↳XGBClassifier(learning_rate=0.05, max_depth=6, use_label_encoder=False) 0.783
xgb = XGBClassifier(n_estimators=100, max_depth=6, learning_rate=0.05,
      ↳use_label_encoder=False)
over = SMOTE(sampling_strategy=0.5, random_state=42)
under = RandomUnderSampler(sampling_strategy=0.35, random_state=42)

steps = [('under', under), ('over', over), ('xgb_clf', xgb)]
smote_model = Pipeline(steps=steps)
```

```
[ ]: smote_model.fit(X_train, y_train)
```

```
[ ]: Pipeline(steps=[('under',
                      RandomUnderSampler(random_state=42, sampling_strategy=0.35)),
                    ('over', SMOTE(random_state=42, sampling_strategy=0.5)),
                    ('xgb_clf',
                     XGBClassifier(learning_rate=0.05, max_depth=6,
                                   use_label_encoder=False))])
```

```
[ ]: y_smote_proba = smote_model.predict_proba(X_test)[: , 1] # using the 'greater
      ↳label', i.e., 0 for 'No default'.
```

```
[ ]: truePR = dict()
falsePR = dict()
roc_auc = dict()

falsePR, truePR, _ = roc_curve(y_test, y_smote_proba)
roc_auc = auc(falsePR, truePR)

plt.style.use('ggplot')
plt.figure(figsize=(16,12))
```


```

lw = 2
plt.plot(
    falsePR,
    truePR,
    color='darkorange',
    lw=lw,
    label="ROC Curve (area = %0.3f)" % roc_auc,
)

plt.plot([0,1], color='navy', lw=lw, linestyle='--', label='Baseline 0.5 - Lack
↳of Skill')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) on
↳XGBClassifier(learning_rate=0.05, max_depth=6, use_label_encoder=False)
↳post-SMOTE Balancing')
plt.legend(loc='lower right')
plt.savefig('./credit/model/best-xgb-model-auc-smote.png')
plt.show()

```

Receiver Operating Characteristic (ROC) on XGBClassifier(learning_rate=0.05, max_depth=6, use_label_encoder=False) post-SMOTE Balancing



```
[ ]: y_smote_preds = smote_model.predict(X_test)
```

```
[ ]: print(classification_report(y_test, y_smote_preds))
```

	precision	recall	f1-score	support
0	0.85	0.93	0.89	4687
1	0.63	0.42	0.50	1313
accuracy			0.82	6000
macro avg	0.74	0.67	0.69	6000
weighted avg	0.80	0.82	0.80	6000

- The **precision** of class 0 : 0.85
- The **recall** (sensitivity) of class 0 : 0.93
- The **false discovery rate** of class 0 : $1 - 0.85 = 0.15$
- The **precision** of class 1, the **minority class** : 0.63
- The **recall** of class 1 : 0.42
- The **false discovery rate** of class 1 : $1 - 0.63 = 0.37$
- The **False Negative Rate** of **Class 1** is $1 - recall$, or $1 - 0.42 = 0.58$ This is **the rate of missing credit defaults**.

```
[ ]: from imblearn.under_sampling import NeighbourhoodCleaningRule
```

```
[ ]: # {'learning_rate': 0.05, 'max_depth': 6, 'n_estimators': 100}
      ↳XGBClassifier(learning_rate=0.05, max_depth=6, use_label_encoder=False) 0.783
xgb = XGBClassifier(n_estimators=100, max_depth=6, learning_rate=0.05,
      ↳use_label_encoder=False)
over = SMOTE(sampling_strategy=0.5, random_state=42)
under = NeighbourhoodCleaningRule(sampling_strategy='auto', n_neighbors=3,
      ↳n_jobs=1)

steps = [('under', under), ('over', over), ('xgb_clf', xgb)]
ncr_model = Pipeline(steps=steps)
```

```
[ ]: ncr_model.fit(X_train, y_train)
```

```
[ ]: Pipeline(steps=[('under', NeighbourhoodCleaningRule(n_jobs=1)),
                    ('over', SMOTE(random_state=42, sampling_strategy=0.5)),
                    ('xgb_clf',
                     XGBClassifier(learning_rate=0.05, max_depth=6,
                                     use_label_encoder=False))])
```

```

[ ]: y_ncr_probas = ncr_model.predict_proba(X_test)[: , 1] # using the 'greater_
      ↪label', i.e., 0 for 'No default'.

[ ]: truePR = dict()
      falsePR = dict()
      roc_auc = dict()


      falsePR, truePR, _ = roc_curve(y_test, y_ncr_probas)
      roc_auc = auc(falsePR, truePR)

      plt.style.use('ggplot')
      plt.figure(figsize=(16,12))
      lw = 2
      plt.plot(
          falsePR,
          truePR,
          color='darkorange',
          lw=lw,
          label="ROC Curve (area = %0.3f)" % roc_auc,
      )

      plt.plot([0,1], color='navy', lw=lw, linestyle='--', label='Baseline 0.5 - Lack_
          ↪of Skill')
      plt.xlim([0.0, 1.0])
      plt.ylim([0.0, 1.05])
      plt.xlabel('False Positive Rate')
      plt.ylabel('True Positive Rate')
      plt.title('Receiver Operating Characteristic (ROC) on_
          ↪XGBClassifier(learning_rate=0.05, max_depth=6, use_label_encoder=False)_
          ↪post-NCR Balancing')
      plt.legend(loc='lower right')
      plt.savefig('./credit/model/best-xgb-model-auc-ncr.png')
      plt.show()

```

Receiver Operating Characteristic (ROC) on XGBClassifier(learning_rate=0.05, max_depth=6, use_label_encoder=False) post-NCR Balancing



```
[ ]: y_ncr_preds = ncr_model.predict(X_test)
```

```
[ ]: print(classification_report(y_test, y_ncr_preds))
```

	precision	recall	f1-score	support
0	0.88	0.86	0.87	4687
1	0.53	0.57	0.55	1313
accuracy			0.79	6000
macro avg	0.70	0.71	0.71	6000
weighted avg	0.80	0.79	0.80	6000

- The **precision** of **class 0** : 0.88;
- The **recall** (sensitivity) of class 0 : 0.86
- The **false discovery rate** of class 0 : $1 - 0.88 = 0.12$
- The **precision** of **class 1**, the **minority class** : 0.53
- The **recall** of class 1 : 0.57
- The **false discovery rate** of class 1 : $1 - 0.53 = 0.47$

[]: